



**Machine**

**Learning**

**Prof. Sergei Gleyzer**

**Lecture**

**PH451, PH551**  
**January 28, 2025**

# Announcements

- **HS #1 due today at 1pm**
- **HS #2 this week, due on Tue. Feb. 4**
- **Textbook: Read Chapters 4, 5**

# Machine Learning

Algorithm choice sets hypothesis Class  $H$

- Goal: find the best function within  $H$ 
  - eg. one that makes the fewest “mistakes”
  - Optimization problem via a learning process
- Evaluate?
  - Loss (Risk) Function on training data
  - Many possible loss functions:
    - Squared
    - Absolute - choice depends on the problem!
    - Cross-entropy

# Regression



Getting Started Prediction Competition

# House Prices - Advanced Regression Techniques

Predict sales prices and practice feature engineering, RFs, and gradient boosting



Kaggle · 4,937 teams · Ongoing

# How Long Will I Live?

Developed by Professors at the University of Pennsylvania and Featured in

TIME

THE WALL STREET JOURNAL

U.S. News & WORLD REPORT


<https://www.blueprintincome.com/tools/life-expectancy-calculator-how-long-will-i-live/>



Gas coma begins to form around nucleus when comet is about 5 AU from Sun.

Nucleus warms and begins to sublimate.

Tail forms, pushed out by solar wind and radiation; distance is now about 1 AU.




Solar heating diminishes; coma and tail disappear between 3 and 5 AU from Sun.



# Regression

**Modify evaluation in induction algorithm**



**Maximum separation**



**Minimal variance**



# Linear Regression

## Linear Regression model prediction

- $y = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_n x_n$  (Eqn. 4-1)

## Vectorized form:

- $y = h_{\theta}(x) = \theta \cdot x$  (Eqn. 4-2)

# Linear Regression

$$y = h_{\theta}(x) = \theta \cdot x$$

$\theta$  is the model's **parameter vector**:

- **bias term**  $\theta_0$  and the **feature weights**  $\theta_1$  to  $\theta_n$

$x$  is the instance's **feature vector**:  $x_0$  to  $x_n$

- $x_0$  **always equal to 1**

$\theta \cdot x$  is the **dot product** of the vectors  $\theta$  and  $x$

$h_{\theta}$  is the **hypothesis function** using the model parameters  $\theta$ .

# Training Linear Regression


**Recall that:**  $\text{RMSE}(\mathbf{X}, h) = \sqrt{\frac{1}{m} \sum_{i=1}^m \left( h(\mathbf{x}^{(i)}) - y^{(i)} \right)^2}$

We can then **minimize**

$$\text{MSE}(\mathbf{X}, h_{\boldsymbol{\theta}}) = \frac{1}{m} \sum_{i=1}^m \left( \boldsymbol{\theta}^\top \mathbf{x}^{(i)} - y^{(i)} \right)^2$$

- How well does the model **fit** the data?

# Least Squares





**Fit by minimizing squares of errors (variance)**


# Gradient Descent

Iterative optimization algorithm to get to an optimal solution or minimize a cost function


- Measure **local gradient** of the error function
- Follow the **steepest gradient** down
  - With each step try to decrease the cost function
  - Until you converge to a minimum



# Gradient Descent




Learning rate  $\alpha$  too small



Learning rate  $\alpha$  too large

# Feature Scaling



Improve learning by making sure all features have **similar scales** (feature scaling)

- **Gradient Descent converges much faster**



# Feature Scaling

## Example:

- Rescale features to have **zero mean** and **unit variance**, i.e.

$$x_j^{(i)} \leftarrow \frac{x_j^{(i)} - \mu_j}{s_j} \quad \text{where } \text{mean } \mu_j = \frac{1}{n} \sum_{i=1}^n x_j^{(i)}$$

and  $s_j$  is **standard deviation** of feature  $j$

# Gradient Descent

## Batch GD


- Use full training dataset (batch)
- Can be slow

## Stochastic GD


- Use random instance
  - Shuffle instances
  - Gradually reduce LR
- Fast, can get out of local minima

## Mini-Batch GD



- Use mini-batches




# Gradient Descent



# Gradient Descent




# Gradient Descent




# Polynomial Regression

- Train powers of features with linear regression
- Watch out for under or overfitting
  - How do you know which?



# Learning Curves

## Validation vs. Training Error:




Underfitting Model



# Learning Curves

## Validation vs. Training Error:



Overfitting Model

# Bias Variance Trade-Off

## Sources of generalization error:

- **Bias:**
  - (**Wrong**) assumptions about data and model
- **Variance:**
  - Model **DoF** and sensitivity to variations in training dataset (many dof, high variance - overfitting model)
- **Resolution** (Irreducible error):
  - Noisy data

**Trade-off!** e.g. **Increase model complexity:**

- **reduce bias, increase variance** and vice versa

# Hands-on Activity