



Machine

Learning

Prof. Sergei Gleyzer

Lecture

PH451, PH551

February 6, 2025

Announcements

- **Hands-on #3/Reading HW #2**
 - due next Tue
 - Textbook: Chapter 6
- **Quiz**
 - Feb 13
- **Mini-hackathon #1 next week**
 - Feb 10-21

Outline

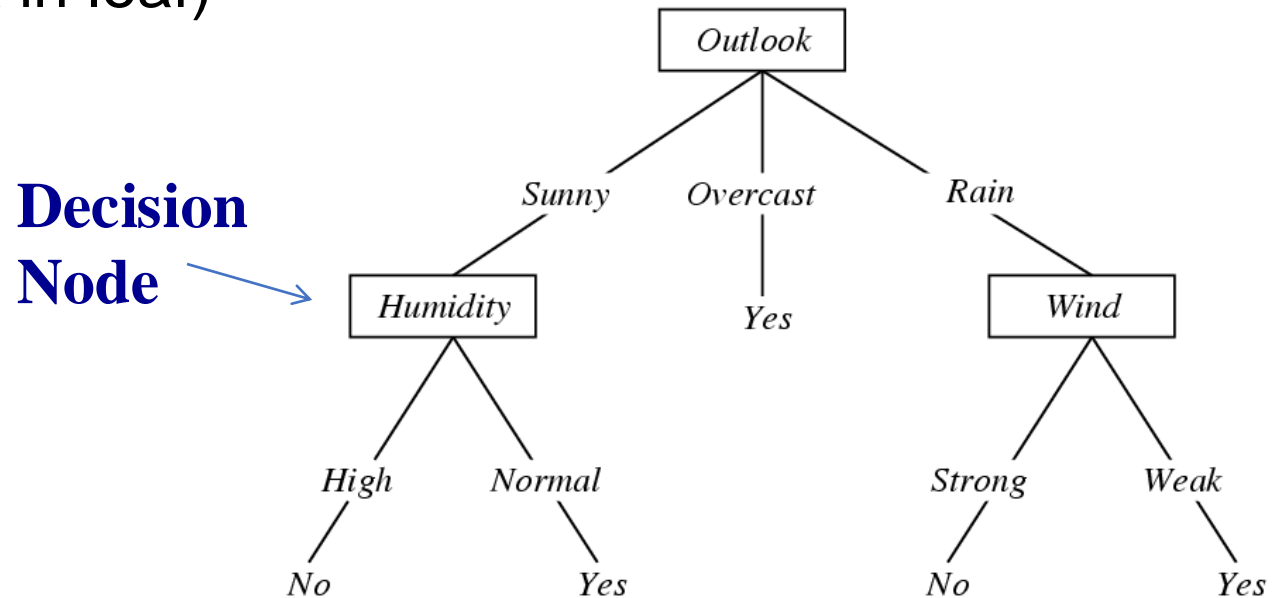
- **Ensembles**
- **Bagging**
- **Random Forests**
- **Boosting**

ML Methods (partial list)

- Fisher, Quadratic
- Naïve Bayes (Likelihood)
- Kernel Density Estimation
- Random Grid Search
- Rule ensembles
- Boosted decision trees
- Random forests
- Deep learning neural networks
- Support vector machines
- Genetic algorithms


Decision Trees (recap)

- Each **internal** node: test one attribute X_i
- Each **branch**: selects one value for X_i
- Each **leaf** node: predict Y
 - Or $P(Y|X \text{ in leaf})$



Decision Trees

- Classic **ML tool** for
 - **decision trees**
 - **rules**
 - **boosted classifiers**
- Written by **J.R. Quinlan**
 - Name: ID3 → C4.5 → C5.0
 - Use c5.0 to familiarize with decision tree classifiers



Pruning

Decision trees can become large and complex and risk **over-fitting** the data

Pruning: remove parts of the tree that are less powerful or possibly **noisy**

- start from the leaves and work back up

Pruned trees smaller in size, easier to interpret

Ensemble Methods



Ensemble Methods

Suppose you have a **collection** of discriminants $f(x, w_k)$, which, individually, perform only **marginally** better than random guessing.

$$f(x) = a_0 + \sum_{k=1}^K a_k f(x, w_k)$$

From such discriminants, **weak learners**, it is possible to build highly effective ones by averaging over them:

Jerome Friedman & Bogdan Popescu (2008)

Ensemble Methods

Bagging (bootstrap aggregation)

- Each tree trained on **bootstrap sample** drawn from training set

Random Forest

- Bagging with **randomized trees**
- Random subsets of features used at each split

Boosting

- Each tree trained on a **different weighting** of full training set. Usually used with decision trees but is more general

Random Forest

Random Forest

- **L. Breinman, 2001**
- **Bagging plus:**
 - Random subset of features for splitting at each node
- **Benefits:**
 - excellent accuracy, avoids over-fitting

Boosting

Turns **weak learners** to **strong learners**
with **weighted ensemble** of **iterative learners**

- Adaptation
- Many boosting algorithms
 - differ in how to weight instances
- Benefits: excellent accuracy
- R. Shapire, 1990

Adaptive Boosting

Adaptive Boosting

Train in stages

- Adaptive weights
 - **ADABOOST: Freund & Schapire 1997**
- **Misclassified** events get a **larger weight** going into the next training stage
 - Classify with a majority vote from all trees
- **Works very well** to improve classification power of “greedy” decision trees

Adaptive Boosting

Repeat K times:

1. Create a decision tree $f(x, w)$
2. Compute its error rate ϵ on the *weighted* training set
3. Compute $\alpha = \ln(1 - \epsilon) / \epsilon$
4. Modify training set: *increase weight* of *incorrectly classified examples* relative to the weights of those that are correctly classified


Then compute weighted average

$$f(x) = \sum \alpha_k f(x, w_k)$$


Y. Freund and R.E. Schapire (1997)

Illustrative Example

$H \rightarrow ZZ^* \rightarrow 4 \text{ leptons}$



$$pp \circledast H \circledast ZZ \circledast \ell^+ \ell^- \ell \ell^+ \ell \ell^-$$









$$pp \circledast ZZ \circledast \ell^+ \ell^- \ell \ell^+ \ell \ell^-$$


$$\mathbf{x} = (m_{Z_1}, m_{Z_2})$$

Credit: H. Prosper







First 6 Decision Trees







First 100 Decision Trees



Averaging over a Forest



H to ZZ to 4Leptons



Build an Ensemble

